

Adversarial Autoencoders, Foundations of Computer Vision, Final Report

Soumyakanti Das

May 2, 2019

1 Introduction

In this report, I have summarized the work that I have done for my final project, which is on Adversarial Autoencoders.

Autoencoders are a type of neural network that has two components, encoder and decoder. Its task is to encode its input into a condensed vector, called Latent Code. The decoder then takes this latent code as input and produces an output which is similar to the input. Thus, it finds application in data compression, novel image generation and dimensionality reduction.

In Adversarial Autoencoders, a Generative Adversarial Network is added to a standard autoencoder, which trains the autoencoder to encode its inputs more efficiently. It does this by imposing a gaussian distribution to the Latent Code space. It should be noted that any arbitrary distribution can also be imposed onto the code space, but since working with gaussian distribution is much easier, I have used it in my project.

I'm using MNIST data for training the network.

2 Architecture

There are three components to the network, Encoder, Decoder, and Discriminator. For the Encoder network, I have three fully connected layers with 25% dropout and ReLU as activation function. The input dimension is the same as the image size, whereas output dimension is just 2. The Decoder network is just a mirror image of the Encoder network. Its input dimension is 2, whereas output dimension is the size of the image. For MNIST, the size of the image tensor is 784 ($28 * 28$). The decoder networks last layer is sigmoidal. The Discriminator network takes two inputs, one from the real gaussian distribution and the other from the output of the encoder. Its input dimension is 2 whereas output dimension is just 1 as it outputs the probability that the code space vector comes from the real distribution. It also has three fully connected layers with 20% dropouts and ReLU activation function.

3 Training

Since there are three networks to be trained simultaneously, the training process is a bit complex. First, in the Reconstruction phase, the encoder and decoder are trained simultaneously to reduce the reconstruction error. This is done by passing the inputs to the encoder, which outputs latent code vector, and then passing the latent code vector to the decoder network. Binary cross entropy loss is calculated on the output of the decoder and the original inputs. This loss is then back propagated to both, encoder and decoder. Next, in the Regularization phase, loss for the discriminator network is calculated on the latent code data and data from actual gaussian distribution. The loss is then backpropagated through the discriminator network. Finally, the generator network, which is also the encoder of the autoencoder, is trained again




Figure 1: Training Losses - MNIST

by calculating loss on its output and the output of the trained discriminator. This loss is then back propagated through the generator (encoder) network.

I have used 0.0005, 0.0001 as my Reconstruction and Regularization learning rates respectively, and used Adam optimizer to training all three networks.

4 Results

After training the networks for 100 epochs, I get the following results.

It can be seen from Figure 1 that Reconstruction and Generator losses decrease because the Encoder (Generator) and Decoder becomes more and more efficient, while Discriminator loss increases, which suggests that it becomes more and more difficult for the Discriminator network to tell apart fake samples from real ones. The encoder learns how to fool the discriminator better with each epoch.

Figure 2 show the generated images after training, from random samples of the gaussian distribution. It can be seen that the decoder is able to generate recognizable images, only from 2 numbers, which is very interesting and shows how efficient the adversarial training is.

Since MNIST and FashionMNIST datasets have images of equal dimensions, I tried to generate images from FashionMNIST with the same architecture. Figure 3 shows the training losses whereas Figure 4 show the generated images after training for 50 epochs for the same architecture.



Figure 2: Generated Images - MNIST




Figure 3: Training Losses - FashionMNIST



Figure 4: Generated Images - FashionMNIST